

Lab 3: Feature Selection - Filter Method

R Abhijit Srivathsan

2448044

Loading California Housing Dataset from sklearn

```
In [13]: import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.feature_selection import VarianceThreshold

# Download the dataset as a pandas DataFrame.
data = fetch_california_housing(as_frame=True)
df = data.frame

# Separate features and target.
X = df.drop('MedHouseVal', axis=1)
y = df['MedHouseVal']

print("Feature names:", X.columns.tolist())
print("Data shape:", X.shape)
```

```
Feature names: ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
Data shape: (20640, 8)
```

In this lab, we load the California housing dataset from sklearn and apply three filter-based feature selection methods:

- **Variance Threshold**
- **Correlation Threshold**
- **Mean Absolute Deviation (MAD)**

Each section shows how adjusting the threshold affects the number of features retained.

Variance Threshold Filtering

The **Variance Threshold** method removes features with low variability. Here we demonstrate:

- The default threshold (0.0) which only removes constant features.
- A higher threshold (0.5) to remove features with variance lower than 0.5.

```
In [14]: # (a) Using the default threshold (0.0) - removes only constant features.
vt_default = VarianceThreshold(threshold=0.0)
X_vt_default = vt_default.fit_transform(X)
print("[VarianceThreshold] With threshold=0.0, shape:", X_vt_default.shape)

# Check the variance of each feature to help decide on a threshold.
variances = X.var()
print("\nVariances of features:\n", variances)

# (b) Increasing the threshold to 0.5.
vt_higher = VarianceThreshold(threshold=0.5)
X_vt_higher = vt_higher.fit_transform(X)
print("\n[VarianceThreshold] With threshold=0.5, shape:", X_vt_higher.shape)

# Get the names of the features that pass the threshold.
selected_features_vt = X.columns[vt_higher.get_support()]
print("Features retained with VarianceThreshold (threshold=0.5):", selected_features_vt.tolist())
```

```
[VarianceThreshold] With threshold=0.0, shape: (20640, 8)
```

Variances of features:

```
MedInc      3.609323e+00
HouseAge    1.583963e+02
AveRooms    6.121533e+00
AveBedrms   2.245915e-01
Population  1.282470e+06
AveOccup    1.078700e+02
Latitude    4.562293e+00
Longitude   4.014139e+00
dtype: float64
```

```
[VarianceThreshold] With threshold=0.5, shape: (20640, 7)
```

```
Features retained with VarianceThreshold (threshold=0.5): ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
```

Correlation Threshold Filtering

The **Correlation Filter** examines pairwise correlations among features. If a pair of features has a correlation above a specified threshold, one of them is dropped.

- We first use a threshold of 0.8.
- Then, we use a stricter threshold of 0.6.

```
In [15]: def correlation_threshold_selector(X, threshold=0.8):
        """
        Identify features to drop based on correlation.

        Parameters:
            X : DataFrame with features.
            threshold : float, correlation coefficient threshold.

        Returns:
            to_drop : list of feature names to drop.
```

```

"""
corr_matrix = X.corr().abs()
# Use only the upper triangle of the correlation matrix.
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))

# Identify columns that have any correlation above the threshold.
to_drop = [column for column in upper.columns if any(upper[column] > threshold)]
return to_drop

# (a) Using a correlation threshold of 0.8.
to_drop_08 = correlation_threshold_selector(X, threshold=0.8)
print("[Correlation Filter] Features dropped (threshold=0.8):", to_drop_08)
selected_features_corr_08 = [col for col in X.columns if col not in to_drop_08]
print("Features retained (correlation threshold=0.8):", selected_features_corr_08)

# (b) Using a stricter threshold of 0.6.
to_drop_06 = correlation_threshold_selector(X, threshold=0.6)
print("\n[Correlation Filter] Features dropped (threshold=0.6):", to_drop_06)
selected_features_corr_06 = [col for col in X.columns if col not in to_drop_06]
print("Features retained (correlation threshold=0.6):", selected_features_corr_06)

```

```

[Correlation Filter] Features dropped (threshold=0.8): ['AveBedrms', 'Longitude']
Features retained (correlation threshold=0.8): ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude']

```

```

[Correlation Filter] Features dropped (threshold=0.6): ['AveBedrms', 'Longitude']
Features retained (correlation threshold=0.6): ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude']

```

Mean Absolute Deviation (MAD) Filtering

The **MAD Filter** removes features with low dispersion (spread). We:

- Calculate the MAD for each feature.
- Select features whose MAD is above a given threshold.
- Experiment with thresholds of 0.5 and 1.0.

```
In [16]: # Calculate the MAD for each feature manually.
mad_values = X.apply(lambda col: np.mean(np.abs(col - col.mean()))))
print("Mean Absolute Deviation (MAD) for each feature:\n", mad_values)

def mad_threshold_selector(X, threshold):
    # Manually calculate MAD for each column.
    mad_values = X.apply(lambda col: np.mean(np.abs(col - col.mean()))))
    selected_features = mad_values[mad_values >= threshold].index.tolist()
    return selected_features

# (a) Applying MAD threshold = 0.5.
selected_features_mad_05 = mad_threshold_selector(X, threshold=0.5)
print("\n[MAD Filter] Features retained (MAD threshold=0.5):", selected_features_mad_05)

# (b) Increasing the threshold to 1.0.
selected_features_mad_10 = mad_threshold_selector(X, threshold=1.0)
print("[MAD Filter] Features retained (MAD threshold=1.0):", selected_features_mad_10)
```

Mean Absolute Deviation (MAD) for each feature:

MedInc	1.401614
HouseAge	10.551539
AveRooms	1.117619
AveBedrms	0.114093
Population	714.237277
AveOccup	0.748172
Latitude	1.975024
Longitude	1.830206

dtype: float64

[MAD Filter] Features retained (MAD threshold=0.5): ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude', 'Longitude']

[MAD Filter] Features retained (MAD threshold=1.0): ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'Latitude', 'Longitude']

Mutual Information for Feature Selection

In this section, we use mutual information to assess the dependency between each feature and the target variable. Mutual information (MI) is a non-negative value that measures the amount of information obtained about one random variable (the target) through another random variable (the feature). A higher MI indicates a more informative feature.

For regression problems, we use `mutual_info_regression`. We then experiment with selecting features that have MI scores above a chosen threshold.

```
In [17]: from sklearn.feature_selection import mutual_info_regression

# Compute the mutual information scores between each feature and the target.
mi_scores = mutual_info_regression(X, y, random_state=0)
mi_series = pd.Series(mi_scores, index=X.columns)

print("Mutual Information Scores (unsorted):")
print(mi_series)

# Sort the scores in descending order for better visualization.
mi_series_sorted = mi_series.sort_values(ascending=False)
print("\nMutual Information Scores (sorted):")
print(mi_series_sorted)
```

Mutual Information Scores (unsorted):

```
MedInc      0.387646
HouseAge    0.032114
AveRooms    0.103623
AveBedrms   0.024482
Population  0.021994
AveOccup    0.072549
Latitude     0.370657
Longitude    0.401820
dtype: float64
```

Mutual Information Scores (sorted):

```
Longitude    0.401820
MedInc       0.387646
Latitude     0.370657
AveRooms     0.103623
AveOccup     0.072549
HouseAge     0.032114
AveBedrms    0.024482
Population   0.021994
dtype: float64
```

```
In [18]: # Define a threshold for selecting features.
# (The threshold value is somewhat arbitrary and may need tuning based on your dataset.)
mi_threshold = 0.1

# Select features that have a mutual information score greater than or equal to the threshold.
selected_features_mi = mi_series[mi_series >= mi_threshold].index.tolist()
print("\nFeatures selected with mutual information score >= {}:".format(mi_threshold))
print(selected_features_mi)
```

```
Features selected with mutual information score >= 0.1:
['MedInc', 'AveRooms', 'Latitude', 'Longitude']
```

Comprehensive Observations on Feature Selection Methods

1. Variance Threshold

- **Threshold = 0.0:**
 - All 8 features are retained since only features with zero variance would be removed.
 - **Dataset Shape:** (20640, 8)
- **Threshold = 0.5:**
 - The feature **AveBedrms** is removed because its variance (~0.2246) is below 0.5.
 - **Retained Features:** ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
 - **Dataset Shape:** (20640, 7)

2. Correlation Filter

- **Thresholds 0.8 and 0.6:**
 - Both thresholds resulted in the same outcome:
 - **Features Dropped:** ['AveBedrms', 'Longitude']
 - **Retained Features:** ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude']
- **Insight:**
 - **AveBedrms** is consistently identified as redundant or uninformative.
 - **Longitude** is removed due to its high correlation with other features—even though, as seen later, it carries significant predictive information according to mutual information analysis.

3. Mean Absolute Deviation (MAD) Filter

- **MAD Values for Each Feature:**

- MedInc : 1.4016
- HouseAge : 10.5515
- AveRooms : 1.1176
- AveBedrms : 0.1141

- Population : 714.2373
- AveOccup : 0.7482
- Latitude : 1.9750
- Longitude : 1.8302
- **Threshold = 0.5:**
 - Removes features with MAD below 0.5, leading to the removal of AveBedrms .
 - **Retained Features:** ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
- **Threshold = 1.0:**
 - With a stricter threshold, features with MAD below 1.0 are removed; here, both AveBedrms and AveOccup (MAD \approx 0.7482) are dropped.
 - **Retained Features:** ['MedInc', 'HouseAge', 'AveRooms', 'Population', 'Latitude', 'Longitude']
- **Insight:**
 - The MAD filter highlights features with very low dispersion. AveBedrms consistently shows minimal variation, while AveOccup is borderline and only gets dropped at higher thresholds.

4. Mutual Information (MI) for Feature Selection

- **Mutual Information Scores (Unsorted):**
 - MedInc : 0.387646
 - HouseAge : 0.032114
 - AveRooms : 0.103623
 - AveBedrms : 0.024482
 - Population : 0.021994
 - AveOccup : 0.072549
 - Latitude : 0.370657
 - Longitude : 0.401820
- **Mutual Information Scores (Sorted Descending):**

1. **Longitude** : 0.401820
2. **MedInc** : 0.387646
3. **Latitude** : 0.370657
4. **AveRooms** : 0.103623
5. **AveOccup** : 0.072549
6. **HouseAge** : 0.032114
7. **AveBedrms** : 0.024482
8. **Population** : 0.021994

- **Feature Selection (Threshold = 0.1):**

- **Selected Features:** ['MedInc', 'AveRooms', 'Latitude', 'Longitude']

- **Insight:**

- The MI analysis reveals that **geographical features** (Longitude and Latitude) and **MedInc** are highly informative.
 - **AveRooms** is marginal but still passes the threshold.
 - Lower MI scores for **HouseAge** , **AveBedrms** , **Population** , and **AveOccup** indicate they contribute less unique predictive information.

Overall Synthesis

- **Consistency:**

- **AveBedrms** is consistently removed or deemed uninformative across all methods (variance, correlation, MAD, and MI), reinforcing its limited value.

- **Complementary Perspectives:**

- **Variance and MAD filters** assess the spread and dispersion of the feature values.
 - The **Correlation filter** targets redundancy between features.
 - **Mutual Information** evaluates how much predictive power each feature contributes toward the target variable.

- **Notable Contrasts:**

- Although the **Correlation filter** removes **Longitude** due to high inter-feature correlation, the **MI analysis** shows that **Longitude** is one of the most informative features, indicating that even highly correlated features can be valuable predictors.

- **Practical Implications:**

- Combining these methods provides a more nuanced view of the data. It helps in balancing dimensionality reduction with the retention of informative features.
- Careful threshold tuning is essential to ensure that the final feature set is both non-redundant and rich in predictive content.