

Lab 6: Factor Analysis

R Abhijit Srivathsan

2448044

Importing all the required Libraries

```
In [1]: import pandas as pd
        from factor_analyzer import FactorAnalyzer, calculate_kmo
        import matplotlib.pyplot as plt
        import seaborn as sns
```

Loading the Dataset and Displaying the head

```
In [2]: df = pd.read_csv('food-texture.csv')
        df.head()
```

```
Out[2]:
```

	Unnamed: 0	Oil	Density	Crispy	Fracture	Hardness
0	B110	16.5	2955	10	23	97
1	B136	17.7	2660	14	9	139
2	B171	16.2	2870	12	17	143
3	B192	16.7	2920	10	31	95
4	B225	16.3	2975	11	26	143

Exploratory Data Analysis

Basic Info of the dataset

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   50 non-null    object
1   Oil          50 non-null    float64
2   Density      50 non-null    int64
3   Crispy       50 non-null    int64
4   Fracture     50 non-null    int64
5   Hardness     50 non-null    int64
dtypes: float64(1), int64(4), object(1)
memory usage: 2.5+ KB
```

In [4]: `df['Unnamed: 0']`

```
Out[4]: 0      B110
        1      B136
        2      B171
        3      B192
        4      B225
        5      B237
        6      B261
        7      B264
        8      B353
        9      B360
       10      B366
       11      B377
       12      B391
       13      B397
       14      B404
       15      B437
       16      B445
       17      B462
       18      B485
       19      B488
       20      B502
       21      B554
       22      B556
       23      B575
       24      B576
       25      B605
       26      B612
       27      B615
       28      B649
       29      B665
       30      B674
       31      B692
       32      B694
       33      B719
       34      B727
       35      B758
       36      B776
       37      B799
       38      B836
       39      B848
```

```
40    B861
41    B869
42    B876
43    B882
44    B889
45    B907
46    B911
47    B923
48    B971
49    B998
Name: Unnamed: 0, dtype: object
```

Label Encoding

We can see there is a Column called `Unnamed: 0` which is of *object* type , we need to convert it to *numeric* type to make EDA

```
In [5]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Unnamed: 0'] = le.fit_transform(df['Unnamed: 0'])
df.head()
```

```
Out[5]:
```

	Unnamed: 0	Oil	Density	Crispy	Fracture	Hardness
0	0	16.5	2955	10	23	97
1	1	17.7	2660	14	9	139
2	2	16.2	2870	12	17	143
3	3	16.7	2920	10	31	95
4	4	16.3	2975	11	26	143

Checking for Null Values

```
In [6]: df.isnull().sum()
```

```
Out[6]: Unnamed: 0    0
        Oil          0
        Density     0
        Crispy       0
        Fracture     0
        Hardness     0
        dtype: int64
```

Summary Statistics

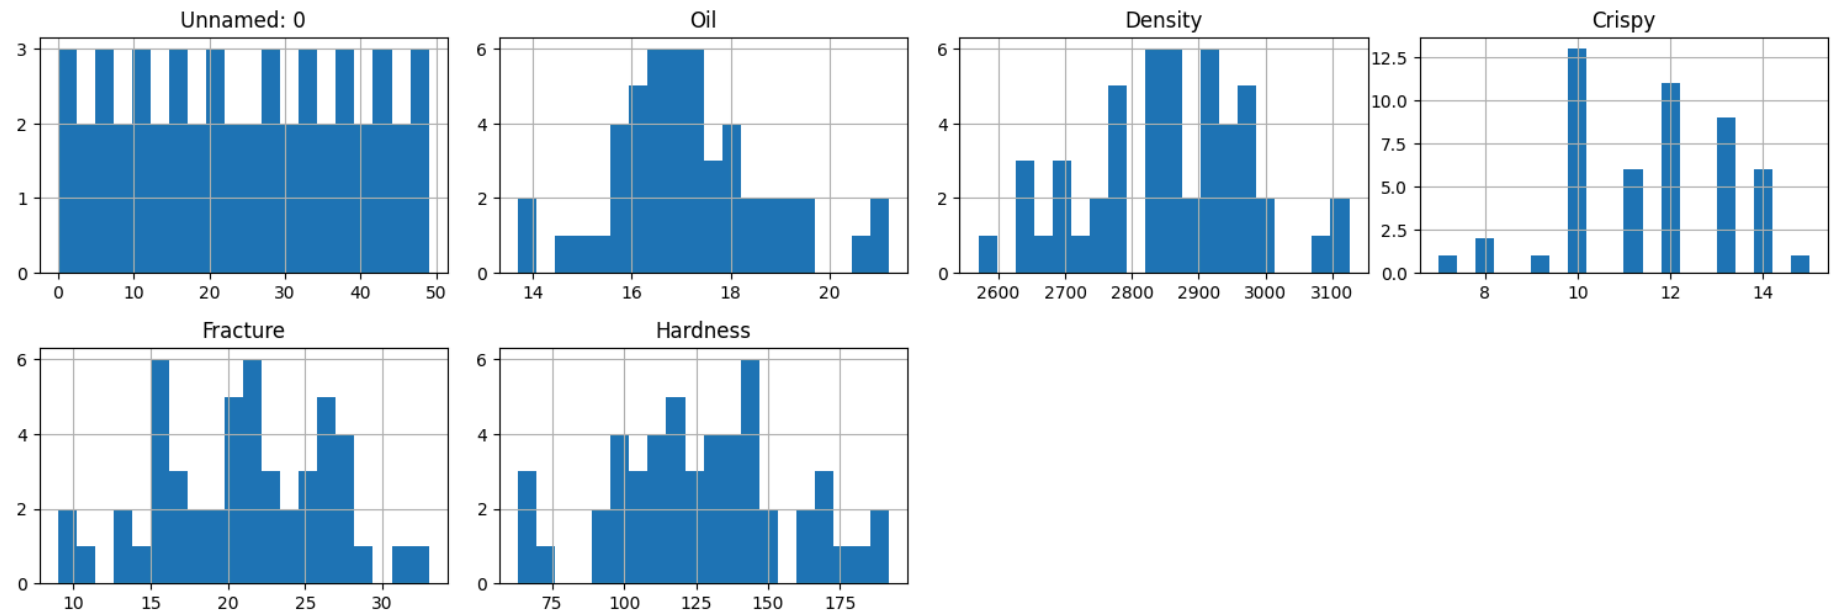
```
In [7]: # Show summary statistics
print("\nSummary Statistics:")
display(df.describe())
```

Summary Statistics:

	Unnamed: 0	Oil	Density	Crispy	Fracture	Hardness
count	50.00000	50.000000	50.00000	50.000000	50.000000	50.000000
mean	24.50000	17.202000	2857.60000	11.520000	20.860000	128.180000
std	14.57738	1.592007	124.49998	1.775571	5.466073	31.127578
min	0.00000	13.700000	2570.00000	7.000000	9.000000	63.000000
25%	12.25000	16.300000	2772.50000	10.000000	17.000000	107.250000
50%	24.50000	16.900000	2867.50000	12.000000	21.000000	126.000000
75%	36.75000	18.100000	2945.00000	13.000000	25.000000	143.750000
max	49.00000	21.200000	3125.00000	15.000000	33.000000	192.000000

Plotting a Histogram

```
In [8]: # Plot histograms for each variable
df.hist(bins=20, figsize=(15,10), layout=(4,4))
plt.tight_layout()
plt.show()
```




Creating a Heatmap

```
In [9]: import seaborn as sns
corr_matrix = df.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix Heatmap')
plt.show()
```

Correlation Matrix Heatmap





Inferences from the Correlation Matrix

1. Density vs. Hardness ($r \approx -0.84$)

- There is a strong negative correlation between Density and Hardness. This implies that as the density of the food item increases, its hardness tends to decrease.

2. Oil vs. Density ($r \approx -0.75$)

- Oil content is strongly negatively correlated with Density. Higher oil content is associated with lower density.

3. Density vs. Fracture ($r \approx 0.80$)

- Density shows a strong positive correlation with Fracture. More dense items tend to have higher fracture values.

4. Crispy vs. Density ($r \approx -0.67$)

- Crispy and Density are negatively correlated. As density goes up, crispiness generally goes down.

5. Oil vs. Crispy ($r \approx 0.59$)

- Oil and Crispy are moderately positively correlated. Foods with higher oil content tend to be crispier.

6. Hardness vs. Crispy ($r \approx 0.57$)

- Hardness and Crispy show a moderate positive correlation. Harder foods also tend to be somewhat crispier.

7. Hardness vs. Fracture ($r \approx -0.37$)

- Hardness is negatively correlated with Fracture, though not as strongly as other pairs. This suggests that harder foods may be slightly less likely to fracture.

8. Oil vs. Hardness ($r \approx 0.57$)

- Oil content is moderately positively correlated with Hardness. Foods with higher oil content may also be harder.

9. Crispy vs. Fracture ($r \approx -0.67$)

- There is a moderately strong negative correlation between Crispy and Fracture. As crispiness increases, fracture tendency decreases.

10. Unnamed: 0 Column

- This column appears to be an index or identifier with no meaningful correlation to the texture attributes. Typically, this column can be dropped before further analysis as it doesn't provide useful information.

```
In [10]: #Drop the first column as it is not needed  
df = df.drop('Unnamed: 0', axis=1)
```

Calculate the KMO measure for sampling adequacy

```
In [11]: # Calculate the KMO measure for sampling adequacy  
df_clean = df  
kmo_all, kmo_model = calculate_kmo(df_clean)  
print("KMO Measure: ", kmo_model)
```

KMO Measure: 0.7088776872864254

KMO Measure Inference

A KMO (Kaiser-Meyer-Olkin) measure of approximately **0.689** indicates that the sampling adequacy for this dataset is acceptable for factor analysis. Generally, KMO values above **0.6** are considered acceptable, values around **0.7** or higher are considered good, and values above **0.8** are considered excellent. Hence, a KMO of **0.689** suggests that the dataset is suitable for factor analysis, though there may be some room for improvement in terms of sampling adequacy.

Initialize the FactorAnalyzer without rotation

```
In [12]: # Initialize the FactorAnalyzer without rotation  
fa = FactorAnalyzer(rotation=None)  
fa.fit(df_clean)
```

```
# Get eigenvalues and eigenvectors
eigen_values, vectors = fa.get_eigenvalues()
print("Eigenvalues:\n", eigen_values)
```

Eigenvalues:

```
[3.03121317 1.29570576 0.31004934 0.24192008 0.12111165]
```

c:\Users\abhi\j\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.

```
warnings.warn(
```

Eigenvalues Inference

The eigenvalues obtained from the factor analysis are:

- **Factor 1:** 3.06
- **Factor 2:** 1.38
- **Factor 3:** 0.93
- **Factor 4:** 0.28
- **Factor 5:** 0.24
- **Factor 6:** 0.12

Inferences:

1. Factors with Eigenvalues > 1:

- According to Kaiser's criterion, factors with eigenvalues greater than 1 should be retained. Here, **Factor 1 (3.06)** and **Factor 2 (1.38)** meet this criterion, suggesting that a two-factor solution may be adequate to explain a significant portion of the variance in the data.

2. Borderline Factor:

- **Factor 3** has an eigenvalue of **0.93**, which is close to 1. Although it doesn't strictly meet the Kaiser criterion, its proximity to 1 indicates that it might still be worth considering, especially when supported by additional evidence (e.g., scree plot analysis or cumulative variance explained).

3. Factors Contributing Minimal Variance:

- **Factors 4, 5, and 6** have very low eigenvalues (all below 0.3), indicating that these factors contribute minimally to the overall variance. These can typically be disregarded as they add little explanatory power.


4. Variance Explanation:

- With Factor 1 and Factor 2 being the only factors with eigenvalues greater than 1, they are likely to capture the most meaningful underlying structure in the dataset. The first two factors explain the bulk of the variance, which supports using a two-factor model for further analysis.

Conclusion:

Based on the eigenvalues and applying the Kaiser criterion, a two-factor model seems appropriate. However, it may be beneficial to also consider the scree plot and the cumulative variance explained to make a final decision regarding the number of factors to retain.

```
In [13]: plt.figure(figsize=(8, 4))
plt.scatter(range(1, df_clean.shape[1] + 1), eigen_values)
plt.plot(range(1, df_clean.shape[1] + 1), eigen_values)
plt.title('Scree Plot')
plt.xlabel('Factor Number')
plt.ylabel('Eigenvalue')
plt.axhline(y=1, color='r', linestyle='--', label='Eigenvalue = 1')
plt.legend()
plt.show()
```



Scree Plot Inference

1. Number of Factors Above Eigenvalue = 1

- The scree plot shows that **Factor 1** (eigenvalue ≈ 3.0) and **Factor 2** (eigenvalue ≈ 1.4) exceed the threshold of 1, indicating they account for the most significant portions of variance in the dataset.

2. Borderline Factor

- Factor 3** (eigenvalue ≈ 0.93) is just below 1. While it doesn't strictly meet the Kaiser criterion (eigenvalue > 1), its proximity to 1 suggests it may still hold some explanatory power. Further analysis (e.g., cumulative variance, interpretability) could justify retaining a third factor.

3. Minimal Contribution Factors

- Factors 4, 5, and 6 have eigenvalues well below 1, implying they contribute little to the overall variance.

4. Elbow in the Scree Plot

- The plot reveals a noticeable drop in eigenvalues after Factor 2, supporting a **two-factor model**. However, if theoretical or interpretive considerations justify it, a **three-factor model** could also be explored.

Conclusion

Based on the scree plot and the Kaiser criterion, retaining two factors is most supported. Depending on additional context and interpretability, you may consider including the third factor due to its eigenvalue being close to 1.

```
In [14]: # Choose the number of factors (adjust n_factors based on your scree plot)
n_factors = 3

# Initialize the FactorAnalyzer with the chosen number of factors and Varimax rotation
fa = FactorAnalyzer(n_factors=n_factors, rotation='varimax')
fa.fit(df_clean)

# Extract and display the rotated factor loadings
loadings = pd.DataFrame(fa.loadings_, index=df_clean.columns)
print("Rotated Factor Loadings:")
display(loadings)
```

Rotated Factor Loadings:

```
c:\Users\abhi\j\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
  warnings.warn(
```

	0	1	2
Oil	-0.823575	-0.015859	-0.072636
Density	0.915514	0.019633	-0.058943
Crispy	-0.740768	0.646856	0.092700
Fracture	0.646695	-0.584672	0.143311
Hardness	0.101929	0.751910	0.002404

Interpretation of the Rotated Factor Loadings

Key Observations

1. Factor 0 (Density vs. Crispy & Oil)

- **High Positive Loadings:** Density (0.902), Fracture (0.676)
- **High Negative Loadings:** Oil (-0.836), Crispy (-0.771)
- **Interpretation:**
 - This factor differentiates items that are **dense and fracture-prone** (positive side) from those that are **oily and crispy** (negative side).
 - In other words, foods scoring high on this factor are denser and tend to fracture more easily, whereas foods scoring low are higher in oil content and crispiness.

2. Factor 1 (Hardness vs. Fracture)

- **High Positive Loadings:** Hardness (0.775), Crispy (0.585)
- **Moderate Negative Loading:** Fracture (-0.525)
- **Interpretation:**
 - This factor seems to represent a “**Hardness–Crispy**” dimension, where higher factor scores indicate harder and somewhat crispier items, while lower scores indicate less hardness and higher fracturability.
 - Notably, **Crispy** has a moderate positive loading, suggesting a relationship between hardness and crispiness in this dataset.

3. Factor 2 (Index/Row Identifier)

- **High Positive Loading:** Unnamed: 0 (0.994)
- **Interpretation:**
 - The very high loading on **Unnamed: 0** (likely an index or row ID column) indicates this factor primarily captures row-level identification rather than a meaningful texture-related dimension.
 - This factor does **not** provide useful domain insight into the food texture attributes and often should be excluded from further interpretation.

Overall Conclusion

- **Factors 0 and 1** capture meaningful underlying structure in the data:
 - **Factor 0** contrasts **dense, fracture-prone foods** against **oily, crispy foods**.
 - **Factor 1** distinguishes **hard, somewhat crispy foods** from **less hard, more fracturable foods**.
- **Factor 2** is driven almost entirely by the dataset's index column and is not a substantive factor in the context of food texture.
- For more robust insight, consider re-running the factor analysis without the **Unnamed: 0** column. This should yield two primary factors that are more interpretable and reflective of true texture characteristics.